

# implementation\_unsupervised

November 30, 2025

## 1 Unsupervised Learning: Implementation Notebook

This notebook explores the dataset without a target label, then applies K-Means to find groups. I kept the steps simple and clearly commented:

- Load features
- Quick EDA
- Scale numeric columns
- PCA for a 2D view
- K-Means clustering
- Silhouette score and a scatter plot

At the end, cluster labels and metrics are saved in `data/processed/`.

```
[1]: # Clean imports for unsupervised workflow
# pandas/numpy for data manipulation; matplotlib/seaborn for visuals.
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
# scikit-learn tooling for scaling, dimensionality reduction, clustering, and evaluation.
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

# Local helper to load the feature-only dataset (no target labels).
from src.data_io import load_unsupervised_features

# Load the dataset. We'll explore and cluster its numeric structure.
unsup_df = load_unsupervised_features()
print("Shape:", unsup_df.shape) # sanity check: rows x columns
# Peek at the first few rows to understand available columns.
unsup_df.head()
```

Shape: (2000, 18)

```
[1]:    age   gender   daily_screen_time_hours   phone_usage_hours  \
0      51        0                  4.8                  3.4
1      64        1                  3.9                  3.5
2      41        2                 10.5                  2.1
3      27        2                  8.8                  0.0
4      55        1                  5.9                  1.7

laptop_usage_hours   tablet_usage_hours   tv_usage_hours   social_media_hours  \
0                  1.3                  1.6                  1.6                  4.1
1                  1.8                  0.9                  2.0                  2.7
2                  2.6                  0.7                  2.2                  3.0
3                  0.0                  0.7                  2.5                  3.3
4                  1.1                  1.5                  1.6                  1.1

work_related_hours   entertainment_hours   gaming_hours  \
0                  2.0                  1.0                  1.7
1                  3.1                  1.0                  1.5
2                  2.8                  4.1                  1.7
3                  1.6                  1.3                  0.4
4                  3.6                  0.8                  0.8

sleep_duration_hours   mood_rating   physical_activity_hours_per_week  \
0                  6.6                  6                      0.7
1                  4.5                  5                      4.3
2                  7.1                  5                      3.1
3                  5.1                 10                     0.0
4                  7.4                  8                      3.0

uses_wellness_apps   eats_healthy   caffeine_intake_mg_per_day  \
0                  1                  1                  125.2
1                  0                  1                  150.4
2                  0                  0                  187.9
3                  0                  1                  73.6
4                  1                  1                  217.5

mindfulness_minutes_per_day
0                  4.0
1                  6.5
2                  6.9
3                  4.8
4                  0.0
```

```
[2]: # Load configuration
from pathlib import Path
import yaml
# Optional config.yaml lets us tweak parameters (e.g., k, output paths) without
# editing code.
```

```

config_path = Path('config.yaml')
config = {}
if config_path.exists():
    with open(config_path, 'r', encoding='utf-8') as f:
        # Safe-load; empty file yields {} so downstream lookups work.
        config = yaml.safe_load(f) or {}
# Echo the active config for transparency in the notebook.
config

```

[2]:

```

{'paths': {'supervised': 'data/raw/clean_supervised.csv',
           'unsupervised': 'data/raw/unsupervised_no_target.csv'},
 'supervised': {'target': 'mental_health_score',
                'test_size': 0.2,
                'random_state': 42},
 'unsupervised': {'use_numeric_only': True, 'random_state': 42, 'k': 3}}

```

## 1.1 Exploratory Data Analysis

[3]:

```

# Quick EDA
# Inspect dtypes and nulls to understand what preprocessing might be needed.
unsup_df.info()
# Summary stats for both numeric/categorical columns provide a fast overview.
display(unsup_df.describe(include='all'))
# Correlation heatmap helps spot collinear numeric features; only meaningful ↴
# with >1 numeric column.
if unsup_df.select_dtypes(include=np.number).shape[1] > 1:
    plt.figure(figsize=(8,6)); sns.heatmap(unsup_df.select_dtypes(include=np.
                                         number).corr(), cmap='coolwarm'); plt.title('Numeric Correlations'); plt.
                                         show()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 18 columns):
 #   Column            Non-Null Count  Dtype  
 ---  --  
 0   age               2000 non-null   int64  
 1   gender             2000 non-null   int64  
 2   daily_screen_time_hours  2000 non-null   float64 
 3   phone_usage_hours   2000 non-null   float64 
 4   laptop_usage_hours  2000 non-null   float64 
 5   tablet_usage_hours  2000 non-null   float64 
 6   tv_usage_hours      2000 non-null   float64 
 7   social_media_hours  2000 non-null   float64 
 8   work_related_hours  2000 non-null   float64 
 9   entertainment_hours 2000 non-null   float64 
 10  gaming_hours        2000 non-null   float64 
 11  sleep_duration_hours 2000 non-null   float64 
 12  mood_rating         2000 non-null   int64  

```

```

13 physical_activity_hours_per_week 2000 non-null float64
14 uses_wellness_apps 2000 non-null int64
15 eats_healthy 2000 non-null int64
16 caffeine_intake_mg_per_day 2000 non-null float64
17 mindfulness_minutes_per_day 2000 non-null float64
dtypes: float64(13), int64(5)
memory usage: 281.4 KB

      age   gender daily_screen_time_hours phone_usage_hours \
count 2000.000000 2000.0000          2000.000000 2000.000000
mean   38.805500  0.6240            6.025600  3.023700
std    14.929203  0.6464            1.974123  1.449399
min   13.000000  0.0000            0.000000  0.000000
25%  26.000000  0.0000            4.700000  2.000000
50%  39.000000  1.0000            6.000000  3.000000
75%  51.000000  1.0000            7.325000  4.000000
max  64.000000  2.0000           13.300000  8.400000

      laptop_usage_hours tablet_usage_hours tv_usage_hours \
count 2000.000000 2000.000000 2000.000000
mean   1.999950  0.995650  1.503700
std    0.997949  0.492714  0.959003
min   0.000000  0.000000  0.000000
25%  1.300000  0.600000  0.800000
50%  2.000000  1.000000  1.500000
75%  2.700000  1.300000  2.200000
max  5.600000  2.500000  4.700000

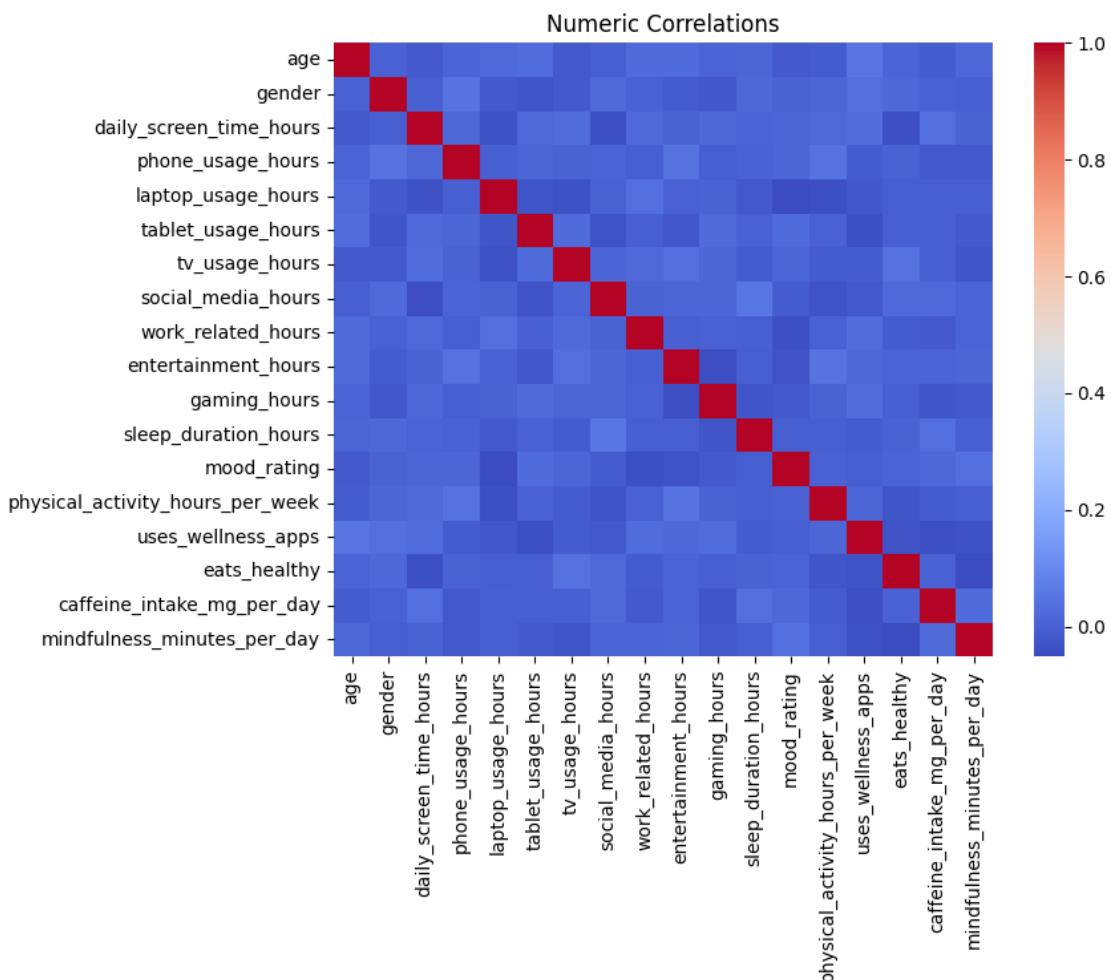
      social_media_hours work_related_hours entertainment_hours \
count 2000.000000 2000.000000 2000.000000
mean   2.039200  2.010250  2.46735
std    1.133435  1.116111  1.23686
min   0.000000  0.000000  0.00000
25%  1.200000  1.200000  1.60000
50%  2.000000  2.000000  2.40000
75%  2.800000  2.800000  3.30000
max  5.800000  5.900000  6.80000

      gaming_hours sleep_duration_hours mood_rating \
count 2000.0000 2000.000000 2000.000000
mean   1.2795  6.537550  5.591000
std    0.8945  1.203856  2.899814
min   0.0000  3.000000  1.000000
25%  0.6000  5.700000  3.000000
50%  1.2000  6.600000  6.000000
75%  1.9000  7.400000  8.000000
max  4.0000 10.000000 10.000000

```

	physical_activity_hours_per_week	uses_wellness_apps	eats_healthy	\
count	2000.000000	2000.000000	2000.000000	
mean	3.087150	0.387500	0.507500	
std	1.885258	0.487301	0.500069	
min	0.000000	0.000000	0.000000	
25%	1.600000	0.000000	0.000000	
50%	3.000000	0.000000	1.000000	
75%	4.400000	1.000000	1.000000	
max	9.700000	1.000000	1.000000	

	caffeine_intake_mg_per_day	mindfulness_minutes_per_day
count	2000.000000	2000.000000
mean	148.07970	10.753750
std	48.86066	7.340269
min	0.80000	0.000000
25%	113.90000	4.900000
50%	147.45000	10.400000
75%	180.70000	15.800000
max	364.90000	36.400000



## 1.2 Scaling, PCA, and Clustering

```
[4]: # Select numeric features (drop non-numeric)
# Unsupervised clustering typically needs numeric inputs; we focus on those.
num_df = unsup_df.select_dtypes(include=np.number).copy()
# Standardize features so PCA/KMeans aren't dominated by scale differences.
scaler = StandardScaler()
X_scaled = scaler.fit_transform(num_df)

# PCA to 2 components for visualization
# Reduces dimensionality; 2D projection lets us plot clusters even if original
# data is high-dimensional.
pca = PCA(n_components=2, random_state=42)
X_pca = pca.fit_transform(X_scaled)
print('Explained variance:', pca.explained_variance_ratio_)

# KMeans clustering (k from config or 3)
# k is the number of clusters; we pull from config to keep it flexible.
k = int(config.get('unsupervised', {}).get('k', 3))
km = KMeans(n_clusters=k, n_init='auto', random_state=42)
labels = km.fit_predict(X_scaled)

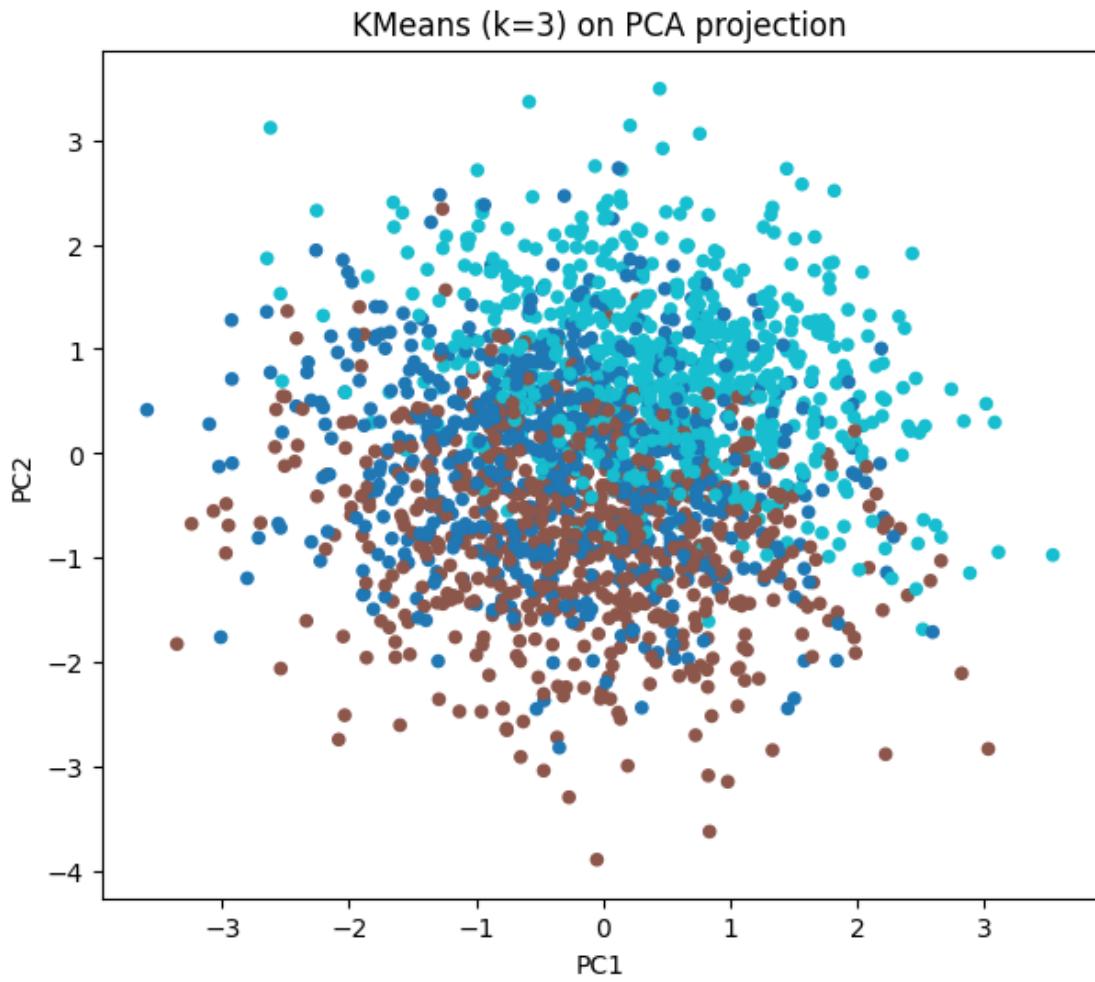
# Silhouette score measures how well points fit within their clusters (higher
# is better).
sil = silhouette_score(X_scaled, labels)
print('Silhouette score:', sil)

# Plot clusters in PCA space for an intuitive visual of separation.
plt.figure(figsize=(7,6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=labels, cmap='tab10', s=20)
plt.title(f'KMeans (k={k}) on PCA projection')
plt.xlabel('PC1'); plt.ylabel('PC2'); plt.show()
```

Explained variance: [0.06426934 0.06384612]

Silhouette score: 0.03798480284443637

Silhouette score: 0.03798480284443637



The silhouette score is very low (~0.038), indicating weak or overlapping cluster structure. This means lifestyle behaviours in this dataset do not naturally separate into clear groups.

### 1.3 Save Artifacts

```
[5]: # Save labels and metrics
from pathlib import Path
import json
# Persist outputs so we can reference them outside the notebook (e.g., in
# slides).
processed_dir = Path(config.get('paths', {}).get('processed', 'data/processed'))
processed_dir.mkdir(parents=True, exist_ok=True)
labels_path = processed_dir / 'unsupervised_kmeans_labels.csv'
metrics_path = processed_dir / 'unsupervised_metrics.json'

# Save cluster assignments as a single-column CSV; easy to join back later.
```

```

out = pd.DataFrame({'cluster': labels})
out.to_csv(labels_path, index=False)
# Record the configuration-derived k and the silhouette score for reporting.
with open(metrics_path, 'w', encoding='utf-8') as f:
    json.dump({'k': int(config.get('unsupervised', {}).get('k', 3)),  

               'silhouette': float(sil)}, f, indent=2)
# Return both paths for quick confirmation.
labels_path, metrics_path

```

[5]: (WindowsPath('data/processed/unsupervised\_kmeans\_labels.csv'),  
 WindowsPath('data/processed/unsupervised\_metrics.json'))

### 1.3.1 K selection via silhouette (2–6)

We test a few k values and report silhouette scores. A low score means clusters overlap and the separation is weak.

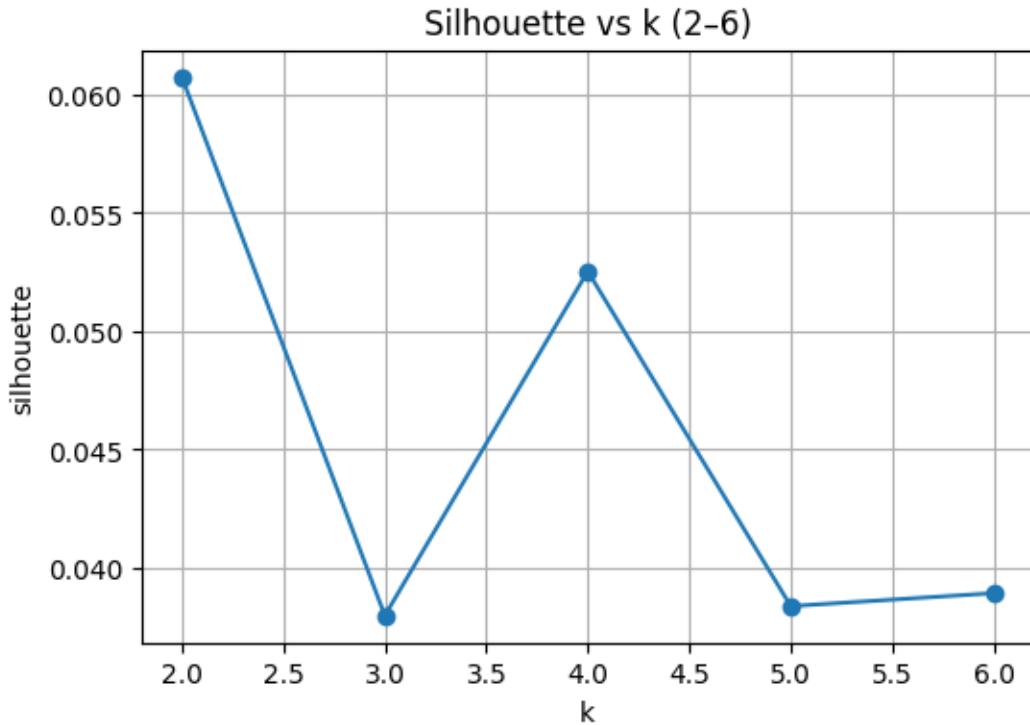
```

[6]: # Silhouette sweep for k=2..6
# We try several k values to see how cluster cohesion/separation changes.
import pandas as pd
from sklearn.metrics import silhouette_score
scores = []
for k_try in range(2, 7):
    km_try = KMeans(n_clusters=k_try, n_init='auto', random_state=42)
    lbl = km_try.fit_predict(X_scaled)
    s = silhouette_score(X_scaled, lbl)
    scores.append({'k': k_try, 'silhouette': float(s)})
# Show a tidy table for quick comparison.
sil_df = pd.DataFrame(scores)
display(sil_df)

# Line plot provides an at-a-glance trend across k.
plt.figure(figsize=(6,4))
plt.plot(sil_df['k'], sil_df['silhouette'], marker='o')
plt.title('Silhouette vs k (2-6)')
plt.xlabel('k')
plt.ylabel('silhouette')
plt.grid(True)
plt.show()

```

k	silhouette
0	0.060755
1	0.037985
2	0.052541
3	0.038384
4	0.038933



### 1.3.2 Cluster interpretation

We join cluster labels back to the original data and compute simple per-cluster means. If the example columns are present, we show those; otherwise we fall back to a few numeric columns.

```
[7]: # Join labels, compute per-cluster means, and print a readable summary
# Attach the computed cluster labels to the original dataframe for ↴interpretation.
unsup_df_with_cluster = unsup_df.copy()
unsup_df_with_cluster['cluster'] = labels

# Prefer these columns if present; else fallback to top 3 numeric columns.
preferred = ['daily_screen_time_hours', 'sleep_duration_hours', 'mood_rating']
available = [c for c in preferred if c in unsup_df_with_cluster.columns]
if not available:
    num_cols = unsup_df_with_cluster.select_dtypes(include=np.number).columns.
    ↴tolist()
    # Avoid including the 'cluster' column in the mean calculation.
    available = [c for c in num_cols if c != 'cluster'][:3]

print('Columns used for interpretation:', available)
# Cluster-wise means give a quick snapshot of typical values per cluster.
```

```

cluster_means = unsup_df_with_cluster.groupby('cluster')[available].mean() .
    ↪round(2)
display(cluster_means)

# Simple textual summary vs overall mean: which features are relatively higher/
    ↪lower per cluster.
overall = unsup_df_with_cluster[available].mean()
print('\nQuick interpretation (relative to overall mean):')
for cl in sorted(unsup_df_with_cluster['cluster'].unique()):
    deltas = (cluster_means.loc[cl] - overall)
    deltas_sorted = deltas.sort_values()
    low = deltas_sorted.index[:1].tolist() # lowest feature relative to overall
    high = deltas_sorted.index[-1:][0].tolist() # highest feature relative to overall
    print(f"- Cluster {cl}: higher {high}; lower {low}")

```

Columns used for interpretation: ['daily\_screen\_time\_hours',  
'sleep\_duration\_hours', 'mood\_rating']

	daily_screen_time_hours	sleep_duration_hours	mood_rating
cluster			
0	6.32	6.65	5.86
1	5.71	6.16	3.73
2	6.05	6.79	7.15

Quick interpretation (relative to overall mean):

- Cluster 0: higher ['daily\_screen\_time\_hours']; lower ['sleep\_duration\_hours']
- Cluster 1: higher ['daily\_screen\_time\_hours']; lower ['mood\_rating']
- Cluster 2: higher ['mood\_rating']; lower ['daily\_screen\_time\_hours']

## 1.4 Cluster Summaries

### 1.4.1 Cluster 0

- Slightly higher screen time
- Slightly less sleep
- Mood around the overall average

Summary: Heavier screen use, a bit less sleep, average mood.

### 1.4.2 Cluster 1

- Lower mood (~3.7 vs ~5.6 overall)
- Sleep and screen time slightly lower than Cluster 0

Summary: Lower mood scores with otherwise average sleep and screen habits.

### 1.4.3 Cluster 2

- Higher mood (~7.1)

- Sleep slightly above average
- Screen time slightly below Cluster 0

Summary: Higher mood and slightly more sleep; somewhat lighter screen use.

Note: Silhouette scores are very low, so boundaries between clusters are fuzzy. Treat these as weak tendencies, not distinct groups.