

implementation_unsupervised

November 30, 2025

1 Unsupervised Learning: Implementation Notebook

This notebook explores the dataset without a target label, then applies K-Means to find groups. I kept the steps simple and clearly commented:

- Load features
- Quick EDA
- Scale numeric columns
- PCA for a 2D view
- K-Means clustering
- Silhouette score and a scatter plot

At the end, cluster labels and metrics are saved in `data/processed/`.

```
[22]: # Clean imports for unsupervised workflow
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score

from src.data_io import load_unsupervised_features

unsup_df = load_unsupervised_features()
print("Shape:", unsup_df.shape)
unsup_df.head()
```

Shape: (2000, 18)

```
[22]:   age  gender  daily_screen_time_hours  phone_usage_hours \
0    51        0                  4.8            3.4
1    64        1                  3.9            3.5
2    41        2                 10.5            2.1
3    27        2                  8.8            0.0
4    55        1                  5.9            1.7
```

	laptop_usage_hours	tablet_usage_hours	tv_usage_hours	social_media_hours	\
0	1.3	1.6	1.6	4.1	
1	1.8	0.9	2.0	2.7	
2	2.6	0.7	2.2	3.0	
3	0.0	0.7	2.5	3.3	
4	1.1	1.5	1.6	1.1	

	work_related_hours	entertainment_hours	gaming_hours	\
0	2.0	1.0	1.7	
1	3.1	1.0	1.5	
2	2.8	4.1	1.7	
3	1.6	1.3	0.4	
4	3.6	0.8	0.8	

	sleep_duration_hours	mood_rating	physical_activity_hours_per_week	\
0	6.6	6	0.7	
1	4.5	5	4.3	
2	7.1	5	3.1	
3	5.1	10	0.0	
4	7.4	8	3.0	

	uses_wellness_apps	eats_healthy	caffeine_intake_mg_per_day	\
0	1	1	125.2	
1	0	1	150.4	
2	0	0	187.9	
3	0	1	73.6	
4	1	1	217.5	

	mindfulness_minutes_per_day
0	4.0
1	6.5
2	6.9
3	4.8
4	0.0

```
[23]: # Load configuration
from pathlib import Path
import yaml
config_path = Path('config.yaml')
config = {}
if config_path.exists():
    with open(config_path, 'r', encoding='utf-8') as f:
        config = yaml.safe_load(f) or {}
config
```

```
[23]: {'paths': {'supervised': 'data/raw/clean_supervised.csv',
               'unsupervised': 'data/raw/unsupervised_no_target.csv'},
```

```
'supervised': {'target': 'mental_health_score',
  'test_size': 0.2,
  'random_state': 42},
'unsupervised': {'use_numeric_only': True, 'random_state': 42, 'k': 3}}
```

1.1 Exploratory Data Analysis

```
[24]: # Quick EDA
unsup_df.info()
display(unsup_df.describe(include='all'))
if unsup_df.select_dtypes(include=np.number).shape[1] > 1:
    plt.figure(figsize=(8,6)); sns.heatmap(unsup_df.select_dtypes(include=np.
    number).corr(), cmap='coolwarm'); plt.title('Numeric Correlations'); plt.
    show()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 18 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   age              2000 non-null   int64  
 1   gender            2000 non-null   int64  
 2   daily_screen_time_hours  2000 non-null   float64 
 3   phone_usage_hours  2000 non-null   float64 
 4   laptop_usage_hours 2000 non-null   float64 
 5   tablet_usage_hours 2000 non-null   float64 
 6   tv_usage_hours    2000 non-null   float64 
 7   social_media_hours 2000 non-null   float64 
 8   work_related_hours 2000 non-null   float64 
 9   entertainment_hours 2000 non-null   float64 
 10  gaming_hours     2000 non-null   float64 
 11  sleep_duration_hours 2000 non-null   float64 
 12  mood_rating      2000 non-null   int64  
 13  physical_activity_hours_per_week 2000 non-null   float64 
 14  uses_wellness_apps 2000 non-null   int64  
 15  eats_healthy      2000 non-null   int64  
 16  caffeine_intake_mg_per_day    2000 non-null   float64 
 17  mindfulness_minutes_per_day  2000 non-null   float64 
dtypes: float64(13), int64(5)
memory usage: 281.4 KB
```

	age	gender	daily_screen_time_hours	phone_usage_hours	\
count	2000.000000	2000.0000	2000.000000	2000.000000	
mean	38.805500	0.6240	6.025600	3.023700	
std	14.929203	0.6464	1.974123	1.449399	
min	13.000000	0.0000	0.000000	0.000000	
25%	26.000000	0.0000	4.700000	2.000000	
50%	39.000000	1.0000	6.000000	3.000000	

75%	51.000000	1.0000	7.325000	4.000000
max	64.000000	2.0000	13.300000	8.400000

	laptop_usage_hours	tablet_usage_hours	tv_usage_hours	\
count	2000.000000	2000.000000	2000.000000	
mean	1.999950	0.995650	1.503700	
std	0.997949	0.492714	0.959003	
min	0.000000	0.000000	0.000000	
25%	1.300000	0.600000	0.800000	
50%	2.000000	1.000000	1.500000	
75%	2.700000	1.300000	2.200000	
max	5.600000	2.500000	4.700000	

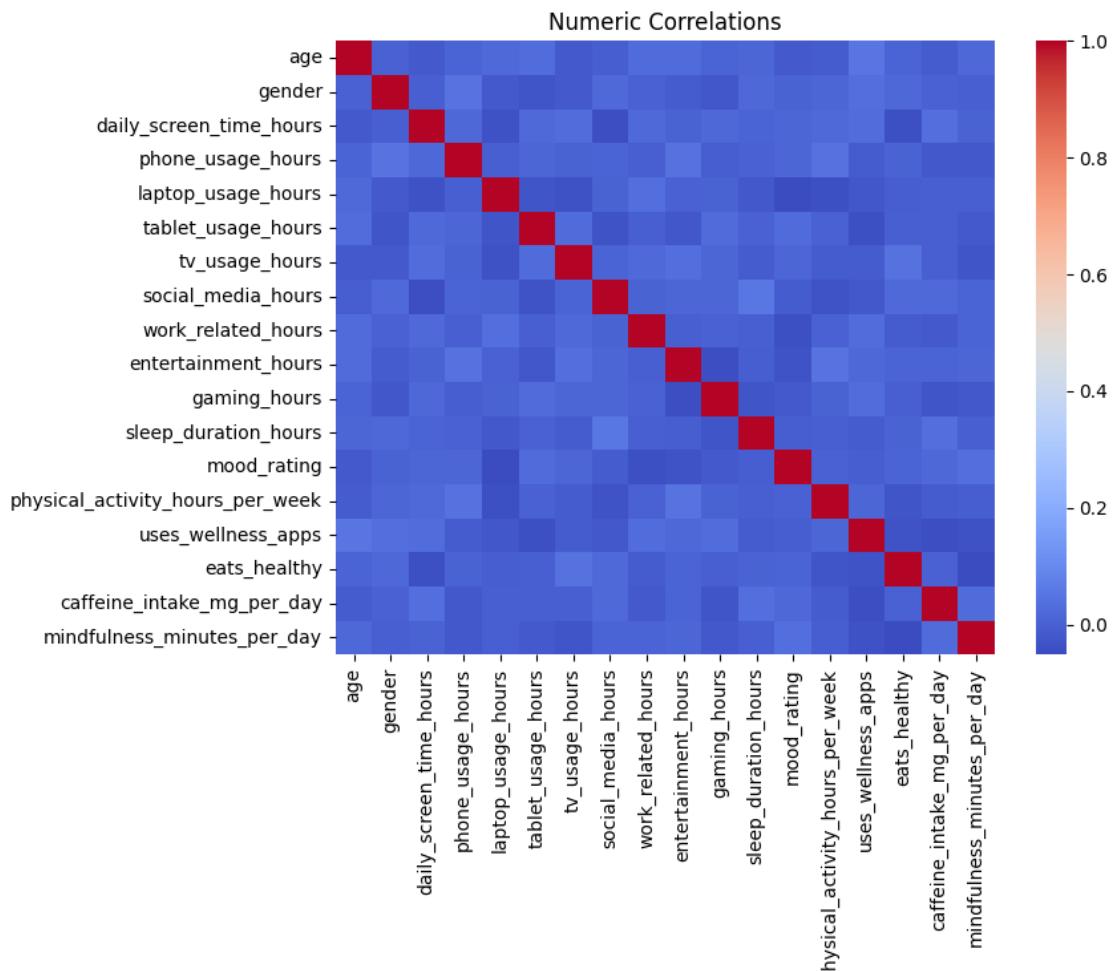
	social_media_hours	work_related_hours	entertainment_hours	\
count	2000.000000	2000.000000	2000.000000	
mean	2.039200	2.010250	2.46735	
std	1.133435	1.116111	1.23686	
min	0.000000	0.000000	0.000000	
25%	1.200000	1.200000	1.60000	
50%	2.000000	2.000000	2.40000	
75%	2.800000	2.800000	3.30000	
max	5.800000	5.900000	6.80000	

	gaming_hours	sleep_duration_hours	mood_rating	\
count	2000.0000	2000.000000	2000.000000	
mean	1.2795	6.537550	5.591000	
std	0.8945	1.203856	2.899814	
min	0.0000	3.000000	1.000000	
25%	0.6000	5.700000	3.000000	
50%	1.2000	6.600000	6.000000	
75%	1.9000	7.400000	8.000000	
max	4.0000	10.000000	10.000000	

	physical_activity_hours_per_week	uses_wellness_apps	eats_healthy	\
count	2000.000000	2000.000000	2000.000000	
mean	3.087150	0.387500	0.507500	
std	1.885258	0.487301	0.500069	
min	0.000000	0.000000	0.000000	
25%	1.600000	0.000000	0.000000	
50%	3.000000	0.000000	1.000000	
75%	4.400000	1.000000	1.000000	
max	9.700000	1.000000	1.000000	

	caffeine_intake_mg_per_day	mindfulness_minutes_per_day	
count	2000.00000	2000.000000	
mean	148.07970	10.753750	
std	48.86066	7.340269	
min	0.80000	0.000000	

25%	113.90000	4.900000
50%	147.45000	10.400000
75%	180.70000	15.800000
max	364.90000	36.400000



1.2 Scaling, PCA, and Clustering

```
[25]: # Select numeric features (drop non-numeric)
num_df = unsup_df.select_dtypes(include=np.number).copy()
scaler = StandardScaler()
X_scaled = scaler.fit_transform(num_df)

# PCA to 2 components for visualization
pca = PCA(n_components=2, random_state=42)
X_pca = pca.fit_transform(X_scaled)
print('Explained variance:', pca.explained_variance_ratio_)
```

```

# KMeans clustering (k from config or 3)
k = int(config.get('unsupervised', {}).get('k', 3))
km = KMeans(n_clusters=k, n_init='auto', random_state=42)
labels = km.fit_predict(X_scaled)

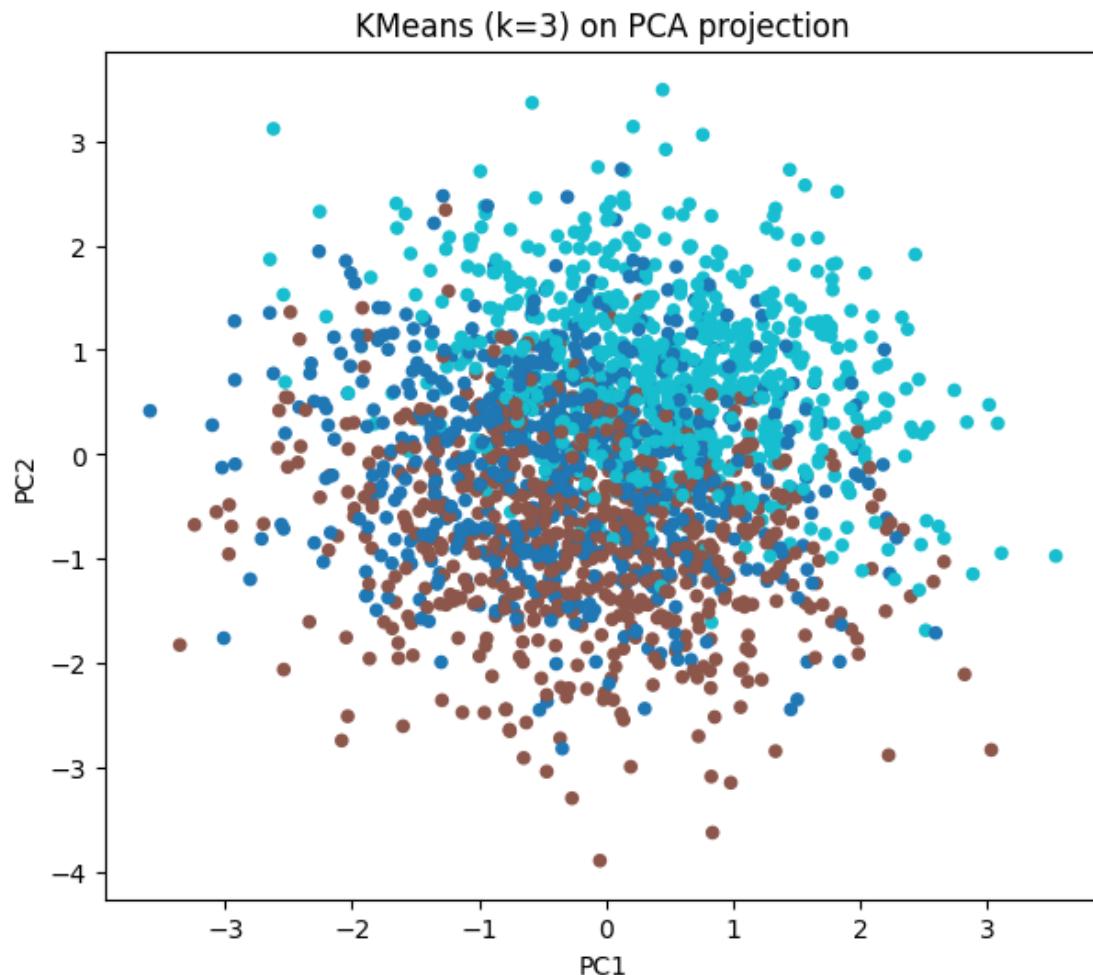
sil = silhouette_score(X_scaled, labels)
print('Silhouette score:', sil)

# Plot clusters in PCA space
plt.figure(figsize=(7,6))
plt.scatter(X_pca[:,0], X_pca[:,1], c=labels, cmap='tab10', s=20)
plt.title(f'KMeans (k={k}) on PCA projection')
plt.xlabel('PC1'); plt.ylabel('PC2'); plt.show()

```

Explained variance: [0.06426934 0.06384612]

Silhouette score: 0.03798480284443637



“The silhouette score is very low (~0.038), indicating weak or overlapping cluster structure. This means lifestyle behaviours in this dataset do not naturally separate into clear groups.”

1.3 Save Artifacts

```
[26]: # Save labels and metrics
from pathlib import Path
import json
processed_dir = Path(config.get('paths', {}).get('processed', 'data/processed'))
processed_dir.mkdir(parents=True, exist_ok=True)
labels_path = processed_dir / 'unsupervised_kmeans_labels.csv'
metrics_path = processed_dir / 'unsupervised_metrics.json'

out = pd.DataFrame({'cluster': labels})
out.to_csv(labels_path, index=False)
with open(metrics_path, 'w', encoding='utf-8') as f:
    json.dump({'k': int(config.get('unsupervised', {}).get('k', 3)), 'silhouette': float(sil)}, f, indent=2)
labels_path, metrics_path
```

```
[26]: (WindowsPath('data/processed/unsupervised_kmeans_labels.csv'),
       WindowsPath('data/processed/unsupervised_metrics.json'))
```

1.3.1 K selection via silhouette (2–6)

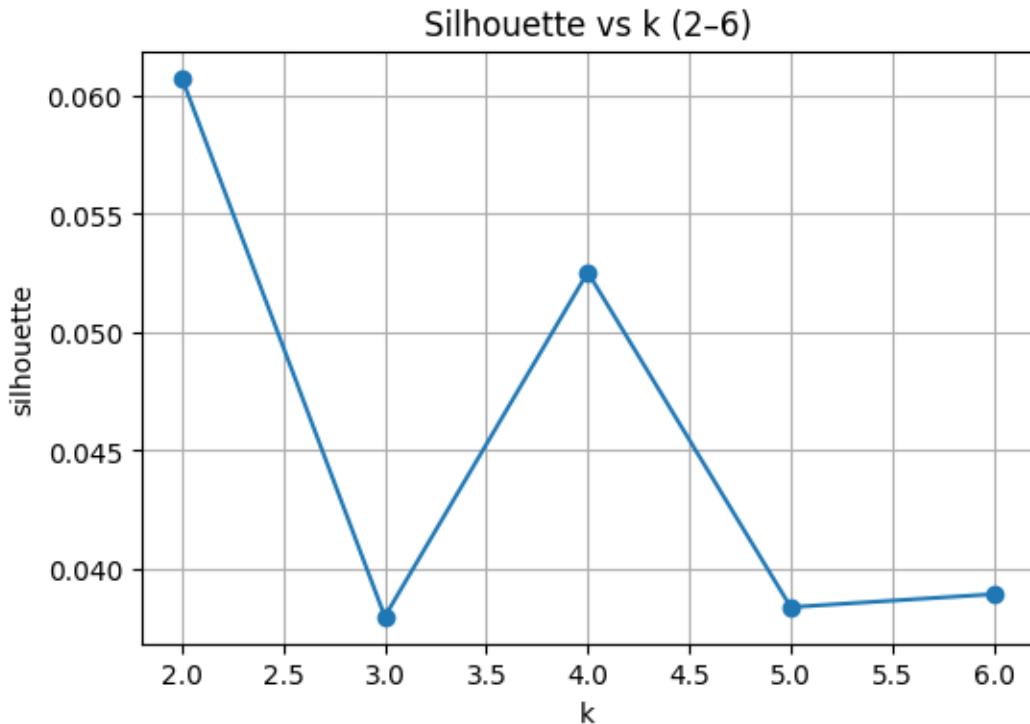
We test a few k values and report silhouette scores. A low score means clusters overlap and the separation is weak.

```
[27]: # Silhouette sweep for k=2..6
import pandas as pd
from sklearn.metrics import silhouette_score
scores = []
for k_try in range(2, 7):
    km_try = KMeans(n_clusters=k_try, n_init='auto', random_state=42)
    lbl = km_try.fit_predict(X_scaled)
    s = silhouette_score(X_scaled, lbl)
    scores.append({'k': k_try, 'silhouette': float(s)})
sil_df = pd.DataFrame(scores)
display(sil_df)

plt.figure(figsize=(6,4))
plt.plot(sil_df['k'], sil_df['silhouette'], marker='o')
plt.title('Silhouette vs k (2-6)')
plt.xlabel('k')
plt.ylabel('silhouette')
plt.grid(True)
```

```
plt.show()
```

```
k    silhouette
0    2    0.060755
1    3    0.037985
2    4    0.052541
3    5    0.038384
4    6    0.038933
```



1.3.2 Cluster interpretation

We join cluster labels back to the original data and compute simple per-cluster means. If the example columns are present, we show those; otherwise we fall back to a few numeric columns.

```
[28]: # Join labels, compute per-cluster means, and print a readable summary
unsup_df_with_cluster = unsup_df.copy()
unsup_df_with_cluster['cluster'] = labels

# Prefer these columns if present; else fallback to top 3 numeric columns
preferred = ['daily_screen_time_hours', 'sleep_duration_hours', 'mood_rating']
available = [c for c in preferred if c in unsup_df_with_cluster.columns]
if not available:
    num_cols = unsup_df_with_cluster.select_dtypes(include=np.number).columns.
    ↴tolist()
```

```

available = [c for c in num_cols if c != 'cluster'][:3]

print('Columns used for interpretation:', available)
cluster_means = unsup_df_with_cluster.groupby('cluster')[available].mean().
    round(2)
display(cluster_means)

# Simple textual summary vs overall mean
overall = unsup_df_with_cluster[available].mean()
print('\nQuick interpretation (relative to overall mean):')
for cl in sorted(unsup_df_with_cluster['cluster'].unique()):
    deltas = (cluster_means.loc[cl] - overall)
    deltas_sorted = deltas.sort_values()
    low = deltas_sorted.index[:1].tolist()
    high = deltas_sorted.index[-1:].tolist()
    print(f"- Cluster {cl}: higher {high}; lower {low}")

```

Columns used for interpretation: ['daily_screen_time_hours',
'sleep_duration_hours', 'mood_rating']

	daily_screen_time_hours	sleep_duration_hours	mood_rating
cluster			
0	6.32	6.65	5.86
1	5.71	6.16	3.73
2	6.05	6.79	7.15

Quick interpretation (relative to overall mean):

- Cluster 0: higher ['daily_screen_time_hours']; lower ['sleep_duration_hours']
- Cluster 1: higher ['daily_screen_time_hours']; lower ['mood_rating']
- Cluster 2: higher ['mood_rating']; lower ['daily_screen_time_hours']